

Kubernetes kubectl Command Study Guide Structure

1. Introduction

- What is Kubernetes and kubectl
- Setting up kubectl
- Kubectl syntax and command structure
- Kubectl autocomplete setup

2. Cluster Management and Configuration

- Checking cluster status
- Managing contexts and configurations
- Switching between clusters and namespaces
- Viewing and managing kubeconfig

3. Basic Resource Operations

- Creating resources
- Viewing and finding resources
- Describing resources in detail
- Deleting resources

4. Deployment Management

- Creating and updating deployments
- Scaling deployments
- Rolling updates and rollbacks
- Checking deployment status

5. Configuration Management

- Working with ConfigMaps
- Managing Secrets
- Setting environment variables

- Resource quotas and limits

6. Pod Operations

- Creating and running pods
- Accessing pod logs
- Executing commands in containers
- Port forwarding to pods

7. Service Management

- Creating and exposing services
- Service discovery
- Managing ingress
- Checking service endpoints

8. Troubleshooting

- Debugging pods and containers
- Viewing logs and events
- Checking resource status
- Common error patterns and solutions

9. Advanced Operations

- Using labels and selectors
- Working with annotations
- Using jsonpath for filtering and formatting
- Patching resources

10. Best Practices

- Namespace organization
- Resource naming conventions
- Using apply vs create
- Declarative vs imperative approaches

Introduction to kubectl

What is Kubernetes and kubectl?

Kubernetes (often abbreviated as K8s) is an open-source container orchestration platform that automates the deployment, scaling, and management of containerized applications. It groups containers that make up an application into logical units for easy management and discovery.

`kubectl` is the official command-line tool for Kubernetes. It allows you to run commands against Kubernetes clusters to: - Deploy applications - Inspect and manage cluster resources - View logs - Execute commands in containers - And much more

As a Kubernetes engineer, `kubectl` will be your primary interface for interacting with Kubernetes clusters on a day-to-day basis.

Setting up kubectl

Before you can use `kubectl`, you need to install it and configure it to communicate with your Kubernetes cluster.

Installation

Here's how to install `kubectl` on different operating systems:

Linux:

```
# Using curl
```

```
curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/  
bin/linux/amd64/kubectl"  
chmod +x kubectl  
sudo mv kubectl /usr/local/bin/
```

```
# Using package manager (Ubuntu/Debian)
```

```
sudo apt-get update && sudo apt-get install -y kubectl
```

```
# Using package manager (CentOS/RHEL)
```

```
cat <<EOF > /etc/yum.repos.d/kubernetes.repo  
[kubernetes]  
name=Kubernetes  
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64  
enabled=1  
gpgcheck=1  
repo_gpgcheck=1
```

```
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg https://  
packages.cloud.google.com/yum/doc/rpm-package-key.gpg  
EOF  
sudo yum install -y kubectl
```

macOS:

```
# Using Homebrew  
brew install kubectl  
  
# Using curl  
curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/  
bin/darwin/amd64/kubectl"  
chmod +x ./kubectl  
sudo mv ./kubectl /usr/local/bin/kubectl
```

Windows:

```
# Using curl  
curl -LO "https://dl.k8s.io/release/v1.26.0/bin/windows/amd64/kubectl.exe"  
# Add the binary to your PATH  
  
# Using Chocolatey  
choco install kubernetes-cli
```

Configuration

After installation, kubectl needs to know which Kubernetes cluster to communicate with and how to authenticate to that cluster. This information is stored in a configuration file called `kubeconfig`.

The default location for the kubeconfig file is `~/.kube/config` on Linux and macOS, and `%USERPROFILE%\kube\config` on Windows.

To check if your kubectl is properly configured:

```
kubectl cluster-info
```

If kubectl is configured correctly, you should see the URL of your Kubernetes master and other services.

Kubectl Syntax and Command Structure

The basic syntax for kubectl commands is:

```
kubectl [command] [TYPE] [NAME] [flags]
```

Where: - **command**: Specifies the operation you want to perform (e.g., create, get, describe, delete) - **TYPE**: Specifies the resource type (e.g., pods, deployments, services) - **NAME**: Specifies the name of the resource (optional for some commands) - **flags**: Specifies optional flags

Common Command Types

1. **Imperative commands**: Direct actions like `run`, `create`, `expose` `bash kubectl create deployment nginx --image=nginx`
2. **Imperative object configuration**: Create objects using configuration files `bash kubectl create -f nginx.yaml`
3. **Declarative object configuration**: Apply configurations to objects `bash kubectl apply -f nginx.yaml`

Basic Command Examples

Get information about the cluster

```
kubectl cluster-info
```

Get a list of all resources in the current namespace

```
kubectl get all
```

Get a list of pods

```
kubectl get pods
```

Get detailed information about a specific pod

```
kubectl describe pod <pod-name>
```

Create a resource from a file

```
kubectl apply -f <filename.yaml>
```

Delete a resource

```
kubectl delete <resource-type> <resource-name>
```

Get logs from a pod

```
kubectl logs <pod-name>
```

```
# Execute a command in a container
```

```
kubectl exec -it <pod-name> -- <command>
```

Kubectl Autocomplete Setup

Setting up autocomplete for kubectl can significantly improve your productivity by reducing typing and helping you discover available commands and options.

Bash Autocomplete

```
# Set up autocomplete in bash
```

```
source <(kubectl completion bash)
```

```
# Add autocomplete permanently to your bash shell
```

```
echo "source <(kubectl completion bash)" >> ~/.bashrc
```

```
# Create a kubectl alias with autocomplete
```

```
echo 'alias k=kubectl' >> ~/.bashrc
```

```
echo 'complete -o default -F __start_kubectl k' >> ~/.bashrc
```

Zsh Autocomplete

```
# Set up autocomplete in zsh
```

```
source <(kubectl completion zsh)
```

```
# Add autocomplete permanently to your zsh shell
```

```
echo '[[ $commands[kubectl] ]] && source <(kubectl completion zsh)' >> ~/.zshrc
```

```
# Create a kubectl alias with autocomplete
```

```
echo 'alias k=kubectl' >> ~/.zshrc
```

```
echo 'compdef __start_kubectl k' >> ~/.zshrc
```

Fish Autocomplete

```
# Set up autocomplete in fish
```

```
kubectl completion fish | source
```

```
# Add autocomplete permanently to your fish shell
```

```
echo 'kubectl completion fish | source' > ~/.config/fish/completions/kubectl.fish
```

Helpful Kubectl Aliases

Creating aliases for commonly used kubectl commands can save you a lot of typing:

```
# Common aliases
alias k='kubectl'
alias kg='kubectl get'
alias kgp='kubectl get pods'
alias kgd='kubectl get deployments'
alias kgs='kubectl get services'
alias kgn='kubectl get nodes'
alias kd='kubectl describe'
alias kdp='kubectl describe pod'
alias kdd='kubectl describe deployment'
alias kds='kubectl describe service'

# Namespace shortcuts
alias kn='kubectl config set-context --current --namespace'
alias kgpa='kubectl get pods --all-namespaces'
alias kgda='kubectl get deployments --all-namespaces'
alias kgsa='kubectl get services --all-namespaces'

# Context management
alias kx='kubectl config use-context'
alias kgc='kubectl config get-contexts'
```

In the next section, we'll dive deeper into cluster management and configuration commands that you'll use regularly as a Kubernetes engineer.

Cluster Management and Configuration

As a Kubernetes engineer, managing cluster connections and configurations is a fundamental skill. This section covers the essential commands for working with clusters, contexts, and configurations.

Checking Cluster Status

Before performing any operations, it's important to verify your connection to the cluster and check its status.

```
# Display cluster information
kubectl cluster-info

# Check the health of cluster components
```

```
kubectl get componentstatuses
```

```
# View all nodes in the cluster
```

```
kubectl get nodes
```

```
# Get detailed information about a specific node
```

```
kubectl describe node <node-name>
```

```
# Check the version of the client and server
```

```
kubectl version
```

```
# Get a short version output
```

```
kubectl version --short
```

Managing Contexts and Configurations

Kubernetes uses contexts to determine which cluster you're communicating with and which user credentials to use. The configuration is stored in the kubeconfig file.

```
# List all available contexts
```

```
kubectl config get-contexts
```

```
# Display the current context
```

```
kubectl config current-context
```

```
# Switch to a different context
```

```
kubectl config use-context <context-name>
```

```
# View the full kubeconfig
```

```
kubectl config view
```

```
# View the full kubeconfig with sensitive data
```

```
kubectl config view --raw
```

Switching Between Clusters and Namespaces

Namespaces provide a way to divide cluster resources between multiple users or projects.

```
# List all namespaces
```

```
kubectl get namespaces
```

```
# Set the default namespace for the current context
```

```
kubectl config set-context --current --namespace=<namespace-name>
```

```
# Run a command in a specific namespace
```



```
kubectl get pods --namespace=<namespace-name>
```

```
# Use the shorthand -n flag for namespace
```

```
kubectl get pods -n <namespace-name>
```

```
# Get resources across all namespaces
```

```
kubectl get pods --all-namespaces
```

```
# Or use the shorthand
```

```
kubectl get pods -A
```

Viewing and Managing Kubeconfig

The kubeconfig file contains clusters, users, and contexts that kubectl uses for communication.

```
# Merge multiple kubeconfig files
```

```
KUBECONFIG=~/.kube/config:~/.kube/another-config kubectl config view --merge
```

```
# Add a new cluster to your kubeconfig
```

```
kubectl config set-cluster <cluster-name> --server=<server-url> --certificate-authority=<ca-file>
```

```
# Add user credentials
```

```
kubectl config set-credentials <user-name> --client-certificate=<cert-file> --client-key=<key-file>
```

```
# Create a new context
```

```
kubectl config set-context <context-name> --cluster=<cluster-name> --user=<user-name> --namespace=<namespace-name>
```

```
# Delete a context from kubeconfig
```

```
kubectl config delete-context <context-name>
```

```
# Delete a cluster from kubeconfig
```

```
kubectl config delete-cluster <cluster-name>
```

```
# Delete a user from kubeconfig
```

```
kubectl config unset users.<user-name>
```

Useful Context Management Aliases

These aliases can help you manage contexts and namespaces more efficiently:

```
# Switch context
```

```
alias kx='kubectl config use-context'
```

```
# Set namespace for current context
alias kn='kubectl config set-context --current --namespace'

# Display current context and namespace
alias kc='kubectl config current-context'
alias kns='kubectl config view --minify | grep namespace | cut -d" " -f6'
```

Basic Resource Operations

This section covers the fundamental operations for working with Kubernetes resources, which you'll use daily as a Kubernetes engineer.

Creating Resources

There are multiple ways to create resources in Kubernetes:

```
# Create a resource from a YAML or JSON file
kubectl create -f <filename.yaml>

# Create a deployment
kubectl create deployment <name> --image=<image>

# Create a service
kubectl create service <type> <name> --tcp=<port>:<target-port>

# Create a namespace
kubectl create namespace <name>

# Create a configmap
kubectl create configmap <name> --from-file=<path> --from-literal=<key>=<value>

# Create a secret
kubectl create secret generic <name> --from-file=<path> --from-literal=<key>=<value>

# Create a job
kubectl create job <name> --image=<image> -- <command>

# Create a cronjob
kubectl create cronjob <name> --image=<image> --schedule="*/1 * * * *" -- <command>
```

Applying Configuration (Declarative Approach)

The `apply` command is the recommended way to manage Kubernetes resources in production:

Apply a configuration file

```
kubectl apply -f <filename.yaml>
```

Apply multiple files

```
kubectl apply -f <file1.yaml> -f <file2.yaml>
```

Apply all files in a directory

```
kubectl apply -f <directory>/
```

Apply configurations from a URL

```
kubectl apply -f https://example.com/manifest.yaml
```

Apply with record flag to record the command in the resource annotation

```
kubectl apply -f <filename.yaml> --record
```

Viewing and Finding Resources

These commands help you list and find resources in your cluster:

List all resources in the current namespace

```
kubectl get all
```

List specific resource types

```
kubectl get pods
```

```
kubectl get deployments
```

```
kubectl get services
```

```
kubectl get configmaps
```

```
kubectl get secrets
```

```
kubectl get nodes
```

```
kubectl get namespaces
```

Get resources with more details

```
kubectl get pods -o wide
```

```
kubectl get deployments -o wide
```

Format output as YAML

```
kubectl get pod <pod-name> -o yaml
```

Format output as JSON

```
kubectl get pod <pod-name> -o json
```

List resources with custom columns

```
kubectl get pods -o custom-  
columns=NAME:.metadata.name,STATUS:.status.phase,NODE:.spec.nodeName
```

Sort resources by a field

```
kubectl get pods --sort-by=.metadata.creationTimestamp
```

Filter resources using a label selector

```
kubectl get pods -l app=nginx
```

List resources across all namespaces

```
kubectl get pods --all-namespaces
```

Describing Resources in Detail

The `describe` command provides detailed information about a resource:

Describe a specific pod

```
kubectl describe pod <pod-name>
```

Describe all pods

```
kubectl describe pods
```

Describe a deployment

```
kubectl describe deployment <deployment-name>
```

Describe a service

```
kubectl describe service <service-name>
```

Describe a node

```
kubectl describe node <node-name>
```

Describe resources selected by label

```
kubectl describe pods -l app=nginx
```

Deleting Resources

Remove resources from your cluster:

Delete a specific resource

```
kubectl delete pod <pod-name>
```

```
kubectl delete deployment <deployment-name>
```

```
kubectl delete service <service-name>
```

Delete using a file

```
kubectl delete -f <filename.yaml>
```

Delete all pods in the current namespace

```
kubectl delete pods --all
```

Delete resources using label selectors

```
kubectl delete pods -l app=nginx
```

Delete a namespace and all its resources

```
kubectl delete namespace <namespace-name>
```

Force delete a pod without waiting for confirmation from the API server

```
kubectl delete pod <pod-name> --grace-period=0 --force
```

Delete all resources in the current namespace

```
kubectl delete all --all
```

Explaining Resources

The `explain` command helps you understand resource definitions:

Get documentation for pod resource

```
kubectl explain pods
```

Get documentation for a specific field

```
kubectl explain pods.spec.containers
```

Get recursive documentation

```
kubectl explain pods --recursive
```

Get documentation for a specific API version

```
kubectl explain pods --api-version=v1
```

Working with Labels and Annotations

Labels and annotations help organize and select your resources:

Add a label to a resource

```
kubectl label pods <pod-name> environment=production
```

Update an existing label

```
kubectl label pods <pod-name> environment=development --overwrite
```

Remove a label

```
kubectl label pods <pod-name> environment-
```

Add an annotation

```
kubectl annotate pods <pod-name> description="My description"
```

Show labels for resources

```
kubectl get pods --show-labels
```

Filter resources by label

```
kubectl get pods -l environment=production
```

```
kubectl get pods -l 'environment in (production,development)'
```

```
kubectl get pods -l 'environment notin (production,development)'
```

These basic operations form the foundation of your daily work with Kubernetes. In the next sections, we'll explore more specific operations for deployments, configurations, and troubleshooting.

Deployment Management

Deployments are one of the most commonly used resources in Kubernetes, allowing you to declaratively manage application updates. This section covers essential commands for working with deployments in your day-to-day operations.

Creating and Updating Deployments

Create a deployment

```
kubectl create deployment nginx --image=nginx
```

Create a deployment with a specific number of replicas

```
kubectl create deployment nginx --image=nginx --replicas=3
```

Create a deployment and expose it as a service

```
kubectl create deployment nginx --image=nginx --port=80
```

Apply a deployment from a YAML file

```
kubectl apply -f deployment.yaml
```

Update a deployment's image

```
kubectl set image deployment/nginx nginx=nginx:1.19
```

Edit a deployment directly

```
kubectl edit deployment nginx
```

Deployment YAML Example

Here's a basic deployment YAML file for reference:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.19
          ports:
            - containerPort: 80
          resources:
            limits:
              cpu: "500m"
              memory: "512Mi"
            requests:
              cpu: "100m"
              memory: "128Mi"
```

Scaling Deployments

Scale a deployment to a specific number of replicas

```
kubectl scale deployment nginx --replicas=5
```

Scale multiple deployments

```
kubectl scale deployment nginx1 nginx2 --replicas=0
```

Scale based on a file

```
kubectl scale --replicas=3 -f deployment.yaml
```

Autoscale a deployment (create a Horizontal Pod Autoscaler)

```
kubectl autoscale deployment nginx --min=2 --max=5 --cpu-percent=80
```

Rolling Updates and Rollbacks

Update a deployment with a new image (triggers a rolling update)
kubectl **set** image deployment/nginx **nginx**=nginx:1.20 --record

Check the status of a rolling update
kubectl rollout status deployment/nginx

Pause a rolling update
kubectl rollout pause deployment/nginx

Resume a paused rolling update
kubectl rollout resume deployment/nginx

View rollout history
kubectl rollout **history** deployment/nginx

View details of a specific revision
kubectl rollout **history** deployment/nginx --revision=2

Rollback to the previous revision
kubectl rollout undo deployment/nginx

Rollback to a specific revision
kubectl rollout undo deployment/nginx --to-revision=2

Restart a deployment (rolling restart of all pods)
kubectl rollout restart deployment/nginx

Checking Deployment Status

Get basic deployment information
kubectl get deployments

Get detailed deployment information
kubectl describe deployment nginx

Check the rollout status
kubectl rollout status deployment/nginx

Get the YAML representation of the current deployment
kubectl get deployment nginx -o yaml

Check ReplicaSets created by the deployment
kubectl get replicaset -l **app**=nginx

Get pods managed by the deployment
kubectl get pods -l **app**=nginx


```
# Watch deployment changes in real-time
kubectl get deployment nginx -w
```

Deployment Strategies

Kubernetes supports different deployment strategies:

1. **RollingUpdate (default):** Gradually replaces old pods with new ones `yaml spec:`
`strategy: type: RollingUpdate rollingUpdate: maxUnavailable: 25% maxSurge: 25%`
2. **Recreate:** Terminates all existing pods before creating new ones `yaml spec:`
`strategy: type: Recreate`

Deployment Troubleshooting

```
# Check deployment events
kubectl describe deployment nginx

# Check ReplicaSet events
kubectl describe rs <replicaset-name>

# Check pod events
kubectl describe pod <pod-name>

# Check pod logs
kubectl logs <pod-name>

# Check previous container logs if a container has restarted
kubectl logs <pod-name> --previous

# Check container status and reasons for crashes
kubectl get pod <pod-name> -o yaml
```

Configuration Management

Configuration management is crucial for maintaining application settings and secrets in Kubernetes. This section covers commands for working with ConfigMaps, Secrets, and other configuration resources.

Working with ConfigMaps

ConfigMaps allow you to decouple configuration from container images.

Create a ConfigMap from literal values

```
kubectl create configmap app-config --from-literal=key1=value1 --from-literal=key2=value2
```

Create a ConfigMap from a file

```
kubectl create configmap app-config --from-file=config.properties
```

Create a ConfigMap from a directory

```
kubectl create configmap app-config --from-file=config-dir/
```

Get ConfigMaps

```
kubectl get configmaps
```

Describe a ConfigMap

```
kubectl describe configmap app-config
```

Edit a ConfigMap

```
kubectl edit configmap app-config
```

Delete a ConfigMap

```
kubectl delete configmap app-config
```

ConfigMap YAML Example

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
data:
  database_url: "mysql://mysql:3306/mydb"
  cache_url: "redis://redis:6379/0"
  ui_properties: |
    color=blue
    size=medium
    shape=round
```

Managing Secrets

Secrets are similar to ConfigMaps but are intended for sensitive data.

Create a Secret from literal values

```
kubectl create secret generic db-credentials --from-literal=username=admin --from-literal=password=secret
```

Create a Secret from files

```
kubectl create secret generic tls-certs --from-file=cert.pem --from-file=key.pem
```

Create a TLS Secret

```
kubectl create secret tls tls-secret --cert=cert.pem --key=key.pem
```

Get Secrets

```
kubectl get secrets
```

Describe a Secret (note that values are not shown)

```
kubectl describe secret db-credentials
```

Get Secret values (base64 encoded)

```
kubectl get secret db-credentials -o yaml
```

Decode a Secret value

```
kubectl get secret db-credentials -o jsonpath='{.data.password}' | base64 --decode
```

Edit a Secret

```
kubectl edit secret db-credentials
```

Delete a Secret

```
kubectl delete secret db-credentials
```

Secret YAML Example

```
apiVersion: v1
kind: Secret
metadata:
  name: db-credentials
type: Opaque
data:
  username: YWRtaW4= # base64 encoded "admin"
  password: c2VjcmV0 # base64 encoded "secret"
```

Setting Environment Variables

Set environment variables in a deployment from a ConfigMap

```
kubectl set env deployment/nginx --from=configmap/app-config
```

Set environment variables in a deployment from a Secret

```
kubectl set env deployment/nginx --from=secret/db-credentials
```

Set a specific environment variable

```
kubectl set env deployment/nginx DB_HOST=mysql
```

Remove an environment variable

```
kubectl set env deployment/nginx DB_HOST-
```

View environment variables in a pod

```
kubectl exec <pod-name> -- env
```

Resource Quotas and Limits

Create a ResourceQuota

```
kubectl create quota my-quota --hard=cpu=1,memory=1G,pods=10
```

Get ResourceQuotas

```
kubectl get resourcequotas
```

Describe a ResourceQuota

```
kubectl describe resourcequota my-quota
```

Create a LimitRange

```
kubectl create -f limit-range.yaml
```

Get LimitRanges

```
kubectl get limitranges
```

Describe a LimitRange

```
kubectl describe limitrange cpu-limit-range
```

LimitRange YAML Example

```
apiVersion: v1
```

```
kind: LimitRange
```

```
metadata:
```

```
  name: cpu-limit-range
```

```
spec:
```

```
  limits:
```

```
  - default:
```

```
    cpu: "1"
```

```
    memory: "512Mi"
```

```
  defaultRequest:
```

```
    cpu: "0.5"
```

```
    memory: "256Mi"
```

```
  type: Container
```

Pod Operations

Pods are the smallest deployable units in Kubernetes. This section covers essential commands for working with pods.

Creating and Running Pods

Create a pod directly (not recommended for production)

```
kubectl run nginx --image=nginx
```

Create a pod with specific resource requests and limits

```
kubectl run nginx --image=nginx --requests=cpu=100m,memory=128Mi --  
limits=cpu=500m,memory=512Mi
```

Create a pod and expose it as a service

```
kubectl run nginx --image=nginx --port=80 --expose
```

Apply a pod from a YAML file

```
kubectl apply -f pod.yaml
```

Pod YAML Example

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx-pod  
  labels:  
    app: nginx  
spec:  
  containers:  
    - name: nginx  
      image: nginx:latest  
      ports:  
        - containerPort: 80  
      resources:  
        limits:  
          cpu: "500m"  
          memory: "512Mi"  
        requests:  
          cpu: "100m"  
          memory: "128Mi"
```

Accessing Pod Logs

Get logs from a pod

```
kubectl logs <pod-name>
```

Get logs from a specific container in a multi-container pod

```
kubectl logs <pod-name> -c <container-name>
```

Stream logs in real-time

```
kubectl logs -f <pod-name>
```

Get logs with timestamps

```
kubectl logs <pod-name> --timestamps
```

Get logs from the previous instance of a container

```
kubectl logs <pod-name> --previous
```

Get a specific number of lines from the log

```
kubectl logs <pod-name> --tail=100
```

Get logs since a specific time

```
kubectl logs <pod-name> --since=1h
```

Executing Commands in Containers

Execute a command in a pod

```
kubectl exec <pod-name> -- ls -la
```

Execute a command in a specific container of a multi-container pod

```
kubectl exec <pod-name> -c <container-name> -- ls -la
```

Start an interactive shell in a pod

```
kubectl exec -it <pod-name> -- /bin/bash
```

Copy files from a pod to local machine

```
kubectl cp <pod-name>:/path/to/file /local/path
```

Copy files from local machine to a pod

```
kubectl cp /local/path <pod-name>:/path/in/pod
```

Port Forwarding to Pods

Forward a local port to a port on the pod

```
kubectl port-forward <pod-name> 8080:80
```

Forward multiple ports

```
kubectrl port-forward <pod-name> 8080:80 8443:443
```

Forward to a deployment, service, or replicaset

```
kubectrl port-forward deployment/nginx 8080:80
```

```
kubectrl port-forward service/nginx 8080:80
```

```
kubectrl port-forward rs/nginx 8080:80
```

Troubleshooting

Effective troubleshooting is essential for maintaining a healthy Kubernetes cluster. This section covers commands and techniques for diagnosing and resolving common issues.

Debugging Pods and Containers

Check pod status

```
kubectrl get pods
```

Get detailed information about a pod

```
kubectrl describe pod <pod-name>
```

Check pod logs

```
kubectrl logs <pod-name>
```

Check previous container logs if a container has restarted

```
kubectrl logs <pod-name> --previous
```

Check events in the namespace

```
kubectrl get events
```

Sort events by timestamp

```
kubectrl get events --sort-by=.metadata.creationTimestamp
```

Create a debugging container

```
kubectrl run debug --image=busybox --rm -it --restart=Never -- sh
```

Create a debugging container in the same namespace as a problematic pod

```
kubectrl run debug --image=busybox --rm -it --restart=Never --  
namespace=<namespace> -- sh
```

Create an ephemeral debug container in an existing pod (Kubernetes v1.18+)

```
kubectrl debug -it <pod-name> --image=busybox --target=<container-name>
```

Viewing Logs and Events

View cluster-wide events

```
kubectl get events --all-namespaces
```

View events for a specific resource

```
kubectl describe pod <pod-name> | grep -A 10 Events:
```

View events for a specific namespace

```
kubectl get events -n <namespace>
```

View warning events only

```
kubectl get events --field-selector type=Warning
```

View events related to a specific object

```
kubectl get events --field-selector involvedObject.name=<pod-name>
```

Stream logs from all containers in a deployment

```
kubectl logs -f deployment/<deployment-name> --all-containers=true
```

Get logs from all pods with a specific label

```
kubectl logs -l app=nginx
```

Checking Resource Status

Check node status

```
kubectl get nodes
```

```
kubectl describe node <node-name>
```

Check pod resource usage

```
kubectl top pods
```

Check node resource usage

```
kubectl top nodes
```

Check pod status with wide output

```
kubectl get pods -o wide
```

Check deployment status

```
kubectl rollout status deployment/<deployment-name>
```

Check service endpoints

```
kubectl get endpoints <service-name>
```

Check if a service is properly connected to pods

```
kubectl describe service <service-name>
```


Common Error Patterns and Solutions

Pod in Pending State

Check if there are resource constraints

```
kubectl describe pod <pod-name> | grep -A 10 Events:
```

Common causes: - Insufficient resources (CPU, memory) - PersistentVolumeClaim not bound - Node selector or affinity rules can't be satisfied

Pod in CrashLoopBackOff State

Check container logs

```
kubectl logs <pod-name>
```

Check previous container logs

```
kubectl logs <pod-name> --previous
```

Common causes: - Application error - Missing configuration - Resource limits too low

Pod in ImagePullBackOff State

Check pod events

```
kubectl describe pod <pod-name> | grep -A 10 Events:
```

Common causes: - Image doesn't exist - Private registry authentication issues - Network connectivity problems

Service Not Routing Traffic

Check service and endpoints

```
kubectl get service <service-name>
```

```
kubectl get endpoints <service-name>
```

Check if pods are running and ready

```
kubectl get pods -l <selector-used-by-service>
```

Check if service selector matches pod labels

```
kubectl describe service <service-name>
```

```
kubectl get pods --show-labels
```

Useful Debugging Commands

Check DNS resolution inside a pod

```
kubectl exec -it <pod-name> -- nslookup kubernetes.default
```

Check network connectivity from a pod

```
kubectl exec -it <pod-name> -- wget -O- http://service-name:port
```

Check API server connectivity

```
kubectl exec -it <pod-name> -- curl -k https://kubernetes.default.svc
```

Get a shell to a node (requires SSH access)

```
ssh <node-name>
```

Check kubelet logs on a node

```
ssh <node-name> 'sudo journalctl -u kubelet'
```

Check container runtime logs

```
ssh <node-name> 'sudo journalctl -u docker' # for Docker
```

```
ssh <node-name> 'sudo journalctl -u containerd' # for containerd
```

These advanced operations will help you manage deployments, handle configurations, work with pods, and troubleshoot issues in your Kubernetes clusters. As you become more familiar with these commands, you'll be able to manage your Kubernetes resources more efficiently and effectively.

Expanded Examples for `kubectl run` Command

The `kubectl run` command is a versatile tool for quickly creating and running pods in your Kubernetes cluster. Here's a comprehensive set of examples to help you understand its capabilities:

Basic Usage

Run a simple nginx pod

```
kubectl run nginx --image=nginx
```

Run a pod with a specific version of an image

```
kubectl run nginx --image=nginx:1.19.0
```

Run a pod and expose it on a specific port

```
kubectl run nginx --image=nginx --port=80
```

Run a pod with a specific label

```
kubectl run nginx --image=nginx --labels="app=web,tier=frontend"
```

Run a pod in a specific namespace

```
kubectl run nginx --image=nginx --namespace=development
```

Resource Management

Run a pod with specific resource requests

```
kubectl run resource-demo --image=nginx --requests=cpu=100m,memory=128Mi
```

Run a pod with both resource requests and limits

```
kubectl run resource-demo --image=nginx --requests=cpu=100m,memory=128Mi --  
limits=cpu=500m,memory=256Mi
```

Run a pod with a specific service account

```
kubectl run sa-demo --image=nginx --serviceaccount=custom-sa
```

Command and Arguments

Run a pod with a specific command

```
kubectl run command-demo --image=busybox --command -- sleep 3600
```

Run a pod with command arguments

```
kubectl run command-demo --image=busybox -- echo "Hello Kubernetes!"
```

Run a pod with multiple command arguments

```
kubectl run command-demo --image=busybox -- /bin/sh -c "while true; do echo  
hello; sleep 10; done"
```

Environment Variables

Run a pod with environment variables

```
kubectl run env-demo --image=nginx --env="DB_HOST=mysql" --  
env="DB_PORT=3306"
```

Run a pod with environment variables from a ConfigMap

```
kubectl run env-demo --image=nginx --env="DB_HOST=$(kubectl get configmap  
db-config -o jsonpath='{.data.host}')
```

Restart Policies

Run a pod with a specific restart policy

```
kubectl run restart-demo --image=busybox --restart=Never -- echo "This pod will not restart"
```

Run a pod that will restart on failure

```
kubectl run restart-demo --image=busybox --restart=OnFailure -- /bin/sh -c "exit 1"
```

Run a pod that will always restart (default behavior)

```
kubectl run restart-demo --image=busybox --restart=Always -- sleep 3600
```

Temporary Pods for Debugging

Run a temporary interactive pod for debugging

```
kubectl run debug --image=busybox --rm -it --restart=Never -- sh
```

Run a temporary pod to test network connectivity

```
kubectl run test-connectivity --image=busybox --rm -it --restart=Never -- wget -O- http://nginx-service:80
```

Run a temporary pod to test DNS resolution

```
kubectl run dns-test --image=busybox --rm -it --restart=Never -- nslookup kubernetes.default
```

Run a temporary pod in a specific namespace for debugging

```
kubectl run debug --image=busybox --rm -it --restart=Never -- namespace=production -- sh
```

Advanced Options

Run a pod with a specific node selector

```
kubectl run node-selector-demo --image=nginx --overrides='{"spec": {"nodeSelector": {"disktype": "ssd"}}}'
```

Run a pod with tolerations

```
kubectl run toleration-demo --image=nginx --overrides='{"spec": {"tolerations": [{"key": "example-key", "operator": "Exists", "effect": "NoSchedule"}]}'
```

Run a pod with a specific priority class

```
kubectl run priority-demo --image=nginx --overrides='{"spec": {"priorityClassName": "high-priority"}}'
```

Run a pod with a specific security context

```
kubectl run security-demo --image=nginx --overrides='{"spec": {"securityContext": {"runAsUser": 1000, "runAsGroup": 3000}}}'
```

Generating YAML Instead of Creating a Pod

Generate YAML for a pod without creating it

```
kubectl run nginx --image=nginx --dry-run=client -o yaml > nginx-pod.yaml
```

Generate YAML with multiple options

```
kubectl run complex-pod --image=nginx --port=80 --labels="app=web" --requests=cpu=100m,memory=128Mi --limits=cpu=500m,memory=256Mi --env="ENV=production" --dry-run=client -o yaml > complex-pod.yaml
```

Combining with Other Commands

Create a pod and immediately get its details

```
kubectl run nginx --image=nginx && kubectl describe pod nginx
```

Create a pod and watch its status

```
kubectl run nginx --image=nginx && kubectl get pod nginx -w
```

Create a pod and immediately get its logs

```
kubectl run log-demo --image=busybox -- echo "Hello logs" && sleep 5 && kubectl logs log-demo
```

Practical Use Cases

Run a database migration job

```
kubectl run db-migration --image=flyway --restart=Never -- flyway migrate
```

Run a data backup job

```
kubectl run backup --image=backup-tool --restart=Never -- backup --source=/data --destination=s3://backup
```

Run a pod for load testing

```
kubectl run load-test --image=locust --restart=Never -- locust -f /tests/locustfile.py --host=http://target-service
```

Run a pod for health checking

```
kubectl run health-check --image=curlimages/curl --restart=Never -- curl -f http://service:8080/health
```

Best Practices

1. **Use Deployments for long-running applications:** While `kubectl run` is convenient for quick tasks, use Deployments for production workloads: `bash kubectl create deployment nginx --image=nginx --replicas=3`
2. **Use Jobs for batch processes:** For completion tasks, use Jobs instead of pods with `restart=Never`: `bash kubectl create job backup --image=backup-tool -- perform-backup`
3. **Use CronJobs for scheduled tasks:** For recurring tasks, use CronJobs: `bash kubectl create cronjob cleanup --image=cleanup-tool --schedule="0 2 * * *" -- cleanup-old-data`
4. **Generate YAML for version control:** Use the `--dry-run=client -o yaml` flags to generate YAML that can be stored in version control: `bash kubectl run nginx --image=nginx --dry-run=client -o yaml > nginx-pod.yaml`
5. **Use labels for organization:** Always add meaningful labels to your pods: `bash kubectl run nginx --image=nginx --labels="app=web,environment=dev,team=frontend"`

These examples demonstrate the flexibility and power of the `kubectl run` command for various scenarios you'll encounter as a Kubernetes engineer.

Expanded Examples for `kubectl create` Command

The `kubectl create` command is essential for creating Kubernetes resources. Here's a comprehensive set of examples to help you understand its capabilities:

Basic Resource Creation

Create a namespace

```
kubectl create namespace development
```

Create a simple pod (though `kubectl run` or `apply` is typically preferred for pods)

```
kubectl create -f pod.yaml
```

Create a deployment

```
kubectl create deployment nginx --image=nginx
```

```
# Create a job
```

```
kubectl create job backup-job --image=backup-tool
```

```
# Create a cronjob
```

```
kubectl create cronjob cleanup --image=cleanup-tool --schedule="0 2 * * *"
```

ConfigMaps and Secrets

```
# Create a ConfigMap from literal values
```

```
kubectl create configmap app-config --from-literal=API_URL=https://  
api.example.com --from-literal=MAX_CONNECTIONS=100
```

```
# Create a ConfigMap from a file
```

```
kubectl create configmap app-config --from-file=config.properties
```

```
# Create a ConfigMap from a directory of files
```

```
kubectl create configmap app-config --from-file=config-dir/
```

```
# Create a ConfigMap from an env file
```

```
kubectl create configmap app-config --from-env-file=.env
```

```
# Create a generic Secret from literal values
```

```
kubectl create secret generic db-credentials --from-literal=username=admin --from-  
literal=password=s3cr3t
```

```
# Create a Secret from files
```

```
kubectl create secret generic tls-certs --from-file=cert.pem --from-file=key.pem
```

```
# Create a TLS Secret
```

```
kubectl create secret tls tls-secret --cert=path/to/tls.cert --key=path/to/tls.key
```

```
# Create a docker-registry Secret for private registry authentication
```

```
kubectl create secret docker-registry regcred --docker-server=<your-registry-  
server> --docker-username=<your-name> --docker-password=<your-password> --  
docker-email=<your-email>
```

Service Resources

```
# Create a ClusterIP service
```

```
kubectl create service clusterip nginx --tcp=80:80
```

```
# Create a NodePort service
```

```
kubectl create service nodeport nginx --tcp=80:80
```

Create a LoadBalancer service

```
kubectl create service loadbalancer nginx --tcp=80:80
```

Create a ExternalName service

```
kubectl create service externalname my-service --external-name example.com
```

RBAC Resources

Create a ServiceAccount

```
kubectl create serviceaccount jenkins
```

Create a Role

```
kubectl create role pod-reader --verb=get,list,watch --resource=pods
```

Create a ClusterRole

```
kubectl create clusterrole pod-reader --verb=get,list,watch --resource=pods
```

Create a RoleBinding

```
kubectl create rolebinding read-pods --role=pod-reader --  
serviceaccount=default:default
```

Create a ClusterRoleBinding

```
kubectl create clusterrolebinding read-pods --clusterrole=pod-reader --  
serviceaccount=default:default
```

Create a RoleBinding for a user

```
kubectl create rolebinding admin --role=admin --user=user1
```

Resource Quotas and Limits

Create a ResourceQuota

```
kubectl create quota my-quota --hard=cpu=1,memory=1G,pods=10
```

Create a ResourceQuota with scopes

```
kubectl create quota my-quota --hard=cpu=1,memory=1G --scopes=BestEffort
```

Generating Resource YAML

Generate a deployment YAML without creating it

```
kubectl create deployment nginx --image=nginx --dry-run=client -o yaml >  
deployment.yaml
```

Generate a service YAML without creating it

```
kubectl create service clusterip nginx --tcp=80:80 --dry-run=client -o yaml >
```


service.yaml

Generate a namespace YAML without creating it

```
kubectl create namespace development --dry-run=client -o yaml > namespace.yaml
```

Generate a configmap YAML without creating it

```
kubectl create configmap app-config --from-literal=API_URL=https://  
api.example.com --dry-run=client -o yaml > configmap.yaml
```

Generate a secret YAML without creating it

```
kubectl create secret generic db-credentials --from-literal=username=admin --from-  
literal=password=s3cr3t --dry-run=client -o yaml > secret.yaml
```

Creating Multiple Resources

Create multiple resources from a directory

```
kubectl create -f ./resources/
```

Create multiple resources from URLs

```
kubectl create -f https://raw.githubusercontent.com/kubernetes/examples/master/  
guestbook/redis-master-deployment.yaml -f https://raw.githubusercontent.com/  
kubernetes/examples/master/guestbook/redis-master-service.yaml
```

Create multiple resources from a combination of files, directories, and URLs

```
kubectl create -f ./deployment.yaml -f ./service/ -f https://example.com/  
configmap.yaml
```

Creating Resources with Additional Options

Create a deployment with a specific number of replicas

```
kubectl create deployment nginx --image=nginx --replicas=3
```

Create a deployment and expose it as a service

```
kubectl create deployment nginx --image=nginx --port=80
```

Create a job with a specific command

```
kubectl create job pi --image=perl -- perl -Mbignum=bpi -wle 'print bpi(2000)'
```

Create a cronjob with specific command arguments

```
kubectl create cronjob date-printer --image=busybox --schedule="*/1 * * * *" -- /  
bin/sh -c "date; echo Hello from Kubernetes"
```

Creating Resources in Different Namespaces

Create a deployment in a specific namespace

```
kubectl create deployment nginx --image=nginx --namespace=development
```

Create a service in a specific namespace

```
kubectl create service clusterip nginx --tcp=80:80 --namespace=development
```

Create a configmap in a specific namespace

```
kubectl create configmap app-config --from-literal=ENV=dev --  
namespace=development
```

Creating Resources with Labels

Create a deployment with labels

```
kubectl create deployment nginx --image=nginx --labels="app=web,tier=frontend"
```

Create a service with labels

```
kubectl create service clusterip nginx --tcp=80:80 --labels="app=web,tier=frontend"
```

Practical Use Cases

Create a complete application stack

```
kubectl create namespace my-app  
kubectl create deployment frontend --image=frontend:v1 --namespace=my-app  
kubectl create deployment backend --image=backend:v1 --namespace=my-app  
kubectl create service clusterip frontend --tcp=80:80 --namespace=my-app  
kubectl create service clusterip backend --tcp=8080:8080 --namespace=my-app  
kubectl create configmap app-config --from-file=config.json --namespace=my-app  
kubectl create secret generic app-secrets --from-file=secrets.json --namespace=my-app
```

Create a database with persistent storage

```
kubectl create -f pvc.yaml  
kubectl create deployment postgres --image=postgres:13 --namespace=database  
kubectl set env deployment/postgres POSTGRES_PASSWORD=secretpassword  
kubectl create service clusterip postgres --tcp=5432:5432 --namespace=database
```

Create a scheduled backup job

```
kubectl create namespace backup  
kubectl create serviceaccount backup-sa --namespace=backup  
kubectl create rolebinding backup-rb --role=backup-role --  
serviceaccount=backup:backup-sa --namespace=backup
```

```
kubectl create cronjob db-backup --image=backup-tool --schedule="0 2 * * *" --  
namespace=backup
```

Best Practices

1. **Use declarative files for production:** While `kubectl create` is convenient for quick tasks, use YAML files with `kubectl apply` for production:
`bash # Generate the YAML first kubectl create deployment nginx --image=nginx --
dry-run=client -o yaml > deployment.yaml # Edit if needed, then apply kubectl
apply -f deployment.yaml`
2. **Use namespaces for organization:** `bash kubectl create namespace development
kubectl create namespace staging kubectl create namespace production`
3. **Add meaningful labels:** `bash kubectl create deployment app --image=app:v1 --
labels="app=myapp,environment=dev,team=frontend"`
4. **Use resource quotas for governance:** `bash kubectl create quota dev-quota --
hard=cpu=4,memory=8G,pods=10 --namespace=development`
5. **Create resources with proper RBAC:** `bash kubectl create serviceaccount app-sa
kubectl create role app-role --verb=get,list,watch --resource=pods,services kubectl
create rolebinding app-rb --role=app-role --serviceaccount=default:app-sa`

These examples demonstrate the versatility of the `kubectl create` command for various scenarios you'll encounter as a Kubernetes engineer.